




1.


- a) Insert the search-key values 31, 41, 40 in order to the following B⁺-tree.



- b) For the resultant B⁺-tree in a), show the form of the tree after deleting 7, 11, 43, 37 and 41 in order.



- c) Re-create the resultant B⁺-tree in a), i.e., rebuild the tree from an empty tree, using *bottom-up B⁺-tree construction*.



2. Consider a hash function on integer search keys i defined by $h(i) = i^2 \bmod B$, where B is the number of buckets. What is wrong with this hash function if $B = 10$?

- Since each integer can be represented as $10a + b$ where $0 \leq b < 10$ and so its square modulo 10 is the same as b^2 modulo 10.
- The squares of 0 to 9 modulo 10 are: 0, 1, 4, 9, 6, 5, 6, 9, 4, 1, and so the buckets for 2,3,7,8 will always be empty and buckets 1, 4, 6, 9 are twice as likely to be hit.
- So, this function is not uniform.

3. A PARTS file with Part# as hash key includes records with the following Part# values: 2369, 3760, 4692, 4871, 5659, 1821, 1074, 7115, 1620, 2428, 3943, 4750, 6975, 4981, 9208. The file uses 8 buckets, numbered 0 to 7. Each bucket is one disk block and holds two records. Load these records into the file in the given order using the hash function $h(K) = K \bmod 8$. Calculate the average number of block accesses for a random retrieval on Part#.

The records will hash to the following buckets:

K	$h(K)$ (bucket number)
2369	1
3760	0
4692	4
4871	7
5659	3
1821	5
1074	2
7115	3
1620	4
2428	4 overflow
3943	7
4750	6
6975	7 overflow
4981	5
9208	0

Two records out of 15 are in overflow buckets, which require an additional block access. The other records require only one block access. Hence, the average time to retrieve a random record is:

$$(1 * (13/15)) + (2 * (2/15)) = 0.867 + 0.266 = 1.133 \text{ block accesses}$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder