

1. Suppose you need to sort a relation r using sort-merge and merge-join the result with an already sorted relation s . Is it possible to pipeline the output of sort-merge of r to the merge join? Explain.

- Assumption: the memory is large enough for the simultaneous execution of the two operations.
- Using pipelining, output from the sort-merge of r is written to a buffer B at the final merge pass.
- When B is full, the merge-join processes tuples from B , joining them with tuples from s until B is empty.
- At this point, the sort-merge operation is resumed and B is refilled.
- This process continues until the merge-join is complete.

2. Consider the query $\Pi_{A,B,C,D} (R \bowtie_{A=C} S)$. Given the following information:

- R is 10 blocks long, and R tuples are 300 bytes long.
 - S is 100 blocks long, and S tuples are 500 bytes long.
 - No common attribute in R and S .
 - The block size is 1024 bytes.
 - Tuples do not span across blocks.
 - Each S tuple joins with exactly one R tuple.
 - The combined size of attributes A , B , C , and D is 450 bytes.
 - A and B are in R and have a combined size of 200 bytes; C and D are in S .
- a) What is the size of the final result, in terms of number of tuples and number of blocks, assuming that the number of duplicates is negligible?
- b) Transform the query so that projection is done before join.
- c) Suppose that three memory blocks are available and *block nested loop join* is used. Suppose that projection (and duplicate elimination) is based on sorting. Compute the cost in terms of number of block transfers for each of the following orders.
- (i) Join followed by projection
- (ii) Projection followed by the join.

a)

- b_R : 10 blocks; $f_R = \lfloor 1024/300 \rfloor = 3$; n_R : 30 tuples
- b_S : 100 blocks; $f_S = \lfloor 1024/500 \rfloor = 2$; n_S : 200 tuples
- Since every S tuple joins with exactly one R tuple, there can be 200 tuples after the join.
- Since the size of the result is 450 bytes/tuple, 2 tuples fit on a block. This means $200 / 2 = 100$ blocks in the final result.

b) By Rule 8(a), we have

$$\Pi_{A,B,C,D} (R \bowtie_{A=C} S) \equiv (\Pi_{A,B}(R)) \bowtie_{A=C} (\Pi_{C,D}(S))$$

c)

(i) Cost of join followed by projection:

- Cost of block nested-loop join = $(10 + 10 \times 100) = \mathbf{1010}$ block transfers
- Number of tuples after join is 200 tuples, each of size 800 bytes. Thus, only one result tuple fits on a block, and **200** blocks have to be written.
- Projection is sort-based using 3 memory blocks. During the initial sorted run creation, unwanted attributes are eliminated on-the-fly to produce tuples of size 450 bytes, i.e., 2 tuples per block. Thus, 200 blocks are scanned and 100 blocks written with $\lceil 100/3 \rceil = 34$ sorted runs created.
- These sorted runs are merged pairwise in $\lceil \log_2 34 \rceil = 6$ passes.
- Projection (and duplicate elimination) cost = $200 + 100 + 100 \times (2 \times 6 - 1) = \mathbf{1400}$ block transfers.
- Total cost = $1010 + 200 + 1400 = 2610$ block transfers.

(ii) Cost of projection followed by join

- Projection is sort-based using 3 memory blocks.
- Sort relation S. During the initial sorted run creation, unwanted attributes are eliminated on-the-fly to produce tuples of size 250 bytes, i.e., 4 tuples per block. Thus, 100 blocks are scanned and 50 blocks written with $\lceil 50/3 \rceil = 17$ sorted runs created.
- These sorted runs are merged pairwise in $\lceil \log_2 17 \rceil = 5$ passes.
- Projection (and duplicate elimination) cost (including writing the result to disk) of S = $100 + 50 + 50 \times (2 \times 5) = \mathbf{650}$ block transfers.
- Sort relation R. During the initial sorted run creation, unwanted attributes are eliminated on-the-fly to produce tuples of size 200 bytes, i.e., 5 tuples per block. Thus, 10 blocks are scanned and 6 blocks written with $\lceil 6/3 \rceil = 2$ sorted runs created.
- These sorted runs are merged pairwise in 1 pass.
- Projection (and duplicate elimination) cost (including writing the result to disk) of R = $10 + 6 + 6 \times 2 = \mathbf{28}$ block transfers.
- The cost of block nested-loop join is $(6 + 6 \times 50) = \mathbf{306}$ block transfers.
- Total cost = $650 + 28 + 306 = 984$ block transfers.